

JBoss Community

Hibernate ORM: Tips, Tricks, and Performance Techniques

Brett Meyer

Senior Software Engineer

Hibernate ORM, Red Hat



HIBERNATE

Brett Meyer

- Hibernate ORM
 - ORM 4 & 5 development
 - Hibernate OSGi
 - Developer community engagement
 - Red Hat enterprise support, Hibernate engineering lead
- Other contributions
 - Apache Camel
 - Infinispan
- Contact me
 - @brettemeyer or +brettmeyer
 - Freenode #hibernate or #hibernate-dev (brmeyer)

[github.com/brmeyer
/HibernateDemos](https://github.com/brmeyer/HibernateDemos)

slideshare.net/brmeyer

ORM? JPA?

- ORM: Object/Relational Mapping
 - Persistence: Data objects outlive the JVM app
 - Maps Java POJOs to relational databases
 - Supports OO concepts: inheritance, object identity, etc.
 - Navigate data by walking the object graph, not the explicit relational model
- JPA: Java Persistence API
- Hibernate ORM provides its own native API, in addition to full JPA support
- Annotations and XML

Overview

- Fetching Strategies
- 2nd Level Entity Cache
- Query Cache
- Cache Management
- Bytecode Enhancement
- Hibernate Search
- Misc. Tips
- **Q&A**

Caveats

- No “one strategy to rule them all”
- Varies greatly between apps
- Important to understand concepts, then apply as necessary
- Does not replace database tuning!

First, the antithesis:

```
public User get(String username) {  
    final Session session = openSession();  
    session.getTransaction().begin();  
    final User user = (User) session.get( User.class, username );  
    session.getTransaction().commit();  
    return user;  
}  
public boolean login(String username) {  
    return get(username) != null;  
}
```

Clean? Yes. But...

EAGER Demo

- Prototypes start “simple”
 - EAGER
 - No caching
 - Overuse of the kitchen sink
- DAO looks clean
- Then you see the SQL logs



Fetching Strategies

Fetching Strategies

- By far, the most frequent mistake
- Also the most costly
- 2 concepts:
 - **WHEN** (fetch timing)
 - **HOW** (fetch style)

WHEN (fetch timing)

Fetch Timing: **EAGER**

- Immediate
- Can be convenient (in some ways)
- Significantly increases payload
- Enormous amount of unnecessary fetches for deep association tree

Fetch Timing: **LAZY**

- Fetched when first accessed
- Collections
 - LAZY by default
 - Utilizes Hibernate's internal concept of “persistent collections”
 - DEMO
- Single attribute (basic)
 - Requires bytecode enhancement
 - Not typically used nor beneficial

Fetch Timing: **LAZY** (cont'd)

- Single (ToOne) associations
 - Fetched when accessed
 - Proxy
 - Default
 - Fetched when accessed (except ID)
 - Handled internally by Hibernate
 - No-proxy
 - Fetched when accessed (including ID)
 - No visible proxy
 - Requires buildtime bytecode enhancement

Fetch Timing: **EXTRA LAZY**

- Collections only
- Fetch individual elements as accessed
- Does not fetch entire collection
- DEMO

HOW (fetch style)

Fetch Style: **JOIN**

- Join fetch (left/outer join)
- Great for ToOne associations
- Multiple collection joins
 - Possible for non-bags
 - Warning: Cartesian product! `SELECT` is normally faster
- DEMO

Fetch Style: **SELECT**

- Follow-up selects
- Default style for collections (avoids cartesian products)
- Can result in 1+N selects (fetch per collection entry)
- DEMO

Fetch Style: **BATCH**

- Fetches multiple entries when one is accessed
- Configurable on both class and collection levels (“batch-size”)
- Simple select and list of keys
- Multiple algorithms (new as of 4.2)
- Determines # of entries to fetch, based on # of provided keys

Fetch Style: **BATCH** (cont'd)

- **Legacy**: pre-determined sizes, rounded down
 - batch-size==25, 24 elements in collection
 - Fetches -> 12, 10, 2
- **Padded**: same as Legacy, size rounded up
- **Dynamic**: builds SQL for the # of keys, limited by “batch-size”
- DEMO

Fetch Style: **SUBSELECT**

- Follow-up select
- Fetches **all** collection entries when accessed for the 1st time
- Original root entry select used as subselect
- Performance depends on the DB itself
- DEMO

Fetch Style: **PROFILE**

- named profile defining fetch styles
- “activated” through the Session

```
@Entity
@FetchProfile(name = "customer-with-orders", fetchOverrides = {
    @FetchProfile.FetchOverride(entity = Customer.class,
        association = "orders", mode = FetchMode.JOIN)
})
public class Customer {
    ...

    @OneToMany
    private Set<Order> orders;
    ...
}
```

```
Session session = ...;
session.enableFetchProfile( "customer-with-orders" );
Customer customer = (Customer) session.get(
    Customer.class, customerId );
```


Fetch Style Tips

- Best to leave defaults in mappings
- Define the style at the query level
- Granular control

2nd Level Entity Cache (2LC)

2LC

- differs from the 1st level (Session)
- SessionFactory level (JVM)
- can scale horizontally on commodity hardware (clustered)
- multiple providers
- currently focus on Infinispan (JBoss) & EHCache (Terracotta)

2LC (cont'd)

- disabled by default
 - can enable globally (not recommended)
 - enable for individual entities and collections
 - configurable at the Session level w/ CacheMode
- cache all properties (default) or only non-lazy

2LC (cont'd)

- Since JVM-level, concurrency is an issue
- Concurrency strategies
 - Read only: simplest and optimal
 - Read-write
 - Supports updates
 - Should not use for serializable transaction isolation
 - Nonstrict-read-write
 - If updates are occasional and unlikely to collide
 - No strict transaction isolation
 - Transactional:
 - Fully transactional cache providers (ie, Infinispan and EHCACHE)
 - JTA required

2LC (cont'd)

- DEMO

Query Cache

Query Cache

- Caches query result sets
- Repetitive queries with identical params
- Different from an entity cache
 - Does not maintain entity state
 - Holds on to sets of identifiers
- If used, best in conjunction with 2LC

Query Cache (cont'd)

- Disabled by default
- Many applications do not gain anything
- Queries invalidated upon relevant updates
- Increases overhead by some: tracks when to invalidate based on commits
- DEMO

Cache Management

Cache Management

- Entities stored in Session cache when saved, updated, or retrieved
- Session#flush syncs all cached with DB
 - Minimize usage
- Manage memory:
 - Session#evict(entity) when no longer needed
 - Session#clear for all
 - Important for large dataset handling
- Same for 2LC, but methods on SessionFactory#getCache

Bytecode Enhancement

Bytecode Enhancement

- Not just for no-proxy, ToOne laziness
- Reduced load on PersistenceContext
 - EntityEntry
 - Internally provides state & Session association
 - “Heavy” operation and typically a hotspot (multiple Maps)
 - ManagedEntity
 - Reduced memory and CPU loads
 - Entities maintain their own state with bytecode enhancement
 - (ManagedEntity)entity.\$\$hibernate_getEntityEntry();
 - 3 ways:
 - Entity implements ManagedEntity and maintains the association
 - Buildtime instrumentation (Ant and recently Gradle/Maven)
 - Runtime instrumentation

Bytecode (cont'd)

- dirtiness checking
 - Legacy: “dirtiness” determined by deep comparison of entity state
 - Enhancement: entity tracks its own dirtiness
 - Simply check the flag (no deep comparison)
 - Already mentioned...
 - Minimize amount of flushes to begin with
 - Minimize the # of entities in the cache -- evict when possible
- Many additional bytecode improvements planned

Hibernate Search

Hibernate Search

- Full-text search on the DB
 - Bad performance
 - CPU/IO overhead
- Offload full-text queries to Hibernate Search engine
 - Fully indexed
 - Horizontally scalable
- Based on Apache Lucene

Hibernate Search (cont'd)

- Annotate entities with @Indexed
- Annotate properties with @Field
 - Index the text: index=Index.YES
 - “Analyze” the text: analyze=Analyze.YES
 - Lucene analyzer
 - Chunks sentences into words
 - Lowercase all of them
 - Exclude common words (“a”, “the”)
- Combo of indexing and analysis == performant full-text searches!

Misc. Tips

Misc. Tips

- Easy to overlook unintended fetches
 - Ex: Fully implemented toString with all associations (fetches everything simply for logging)
- Use @Immutable when possible
 - Excludes entity from dirtiness checks
 - Other internal optimizations

Misc. Tips (cont'd)

- Use Bag or List for inverse collections, not Set
 - Collection#add & #addAll always return true (don't need to check for pre-existing values)
 - I.e., can add a value without fetching the collection

```
Parent p = (Parent) session.load(Parent.class, id);
Child c = new Child();
c.setParent(p);
p.addChild(c);
...
// does *not* fetch the collection
p.getChildren().add(c);
```

Misc. Tips (cont'd)

- One-shot delete (non-inverse collections)
 - Scenario: 20 elements, need to delete 18 & add 3
 - Default: Hibernate would delete the 18 one-by-one, then add the 3
 - Instead:
 - Discard the entire collection (deletes all elements)
 - Save a new collection with 5 elements
 - 1 bulk delete SQL and 5 inserts
 - Important concept for large amounts of data

How to Help:

hibernate.org
/orm/contribute

QUESTIONS?

- Q&A
- #hibernate or #hibernate-dev (brmeyer)
- @brettemeyer
- +brettmeyer