

# INTRODUCTION TO ELASTICSEARCH

---

# Agenda

- Me
- ElasticSearch Basics
  - Concepts
  - Network / Discovery
  - Data Structure
  - Inverted Index
- The REST API
- Sample Deployment

# Me

- Roy Russo
- JBoss Portal Co-Founder
- LoopFuse Co-Founder
- ElasticHQ
  - <http://www.elastichq.org>
- AltiSource Labs Architect



# ElasticSearch in One Slide

- Document - Oriented Search Engine
  - JSON
  - Apache Lucene
- No Schema
  - Mapping Types
- Horizontal Scale, Distributed
- REST API
- Vibrant Ecosystem
  - Tooling, Plugins, Hosting, Client-Libs

# When to use Elasticsearch


- Full-Text Search
- Fast Read Database
- “Simple” Data Structures
- Minimize Impedance Mismatch

# When to use Elasticsearch - Logs

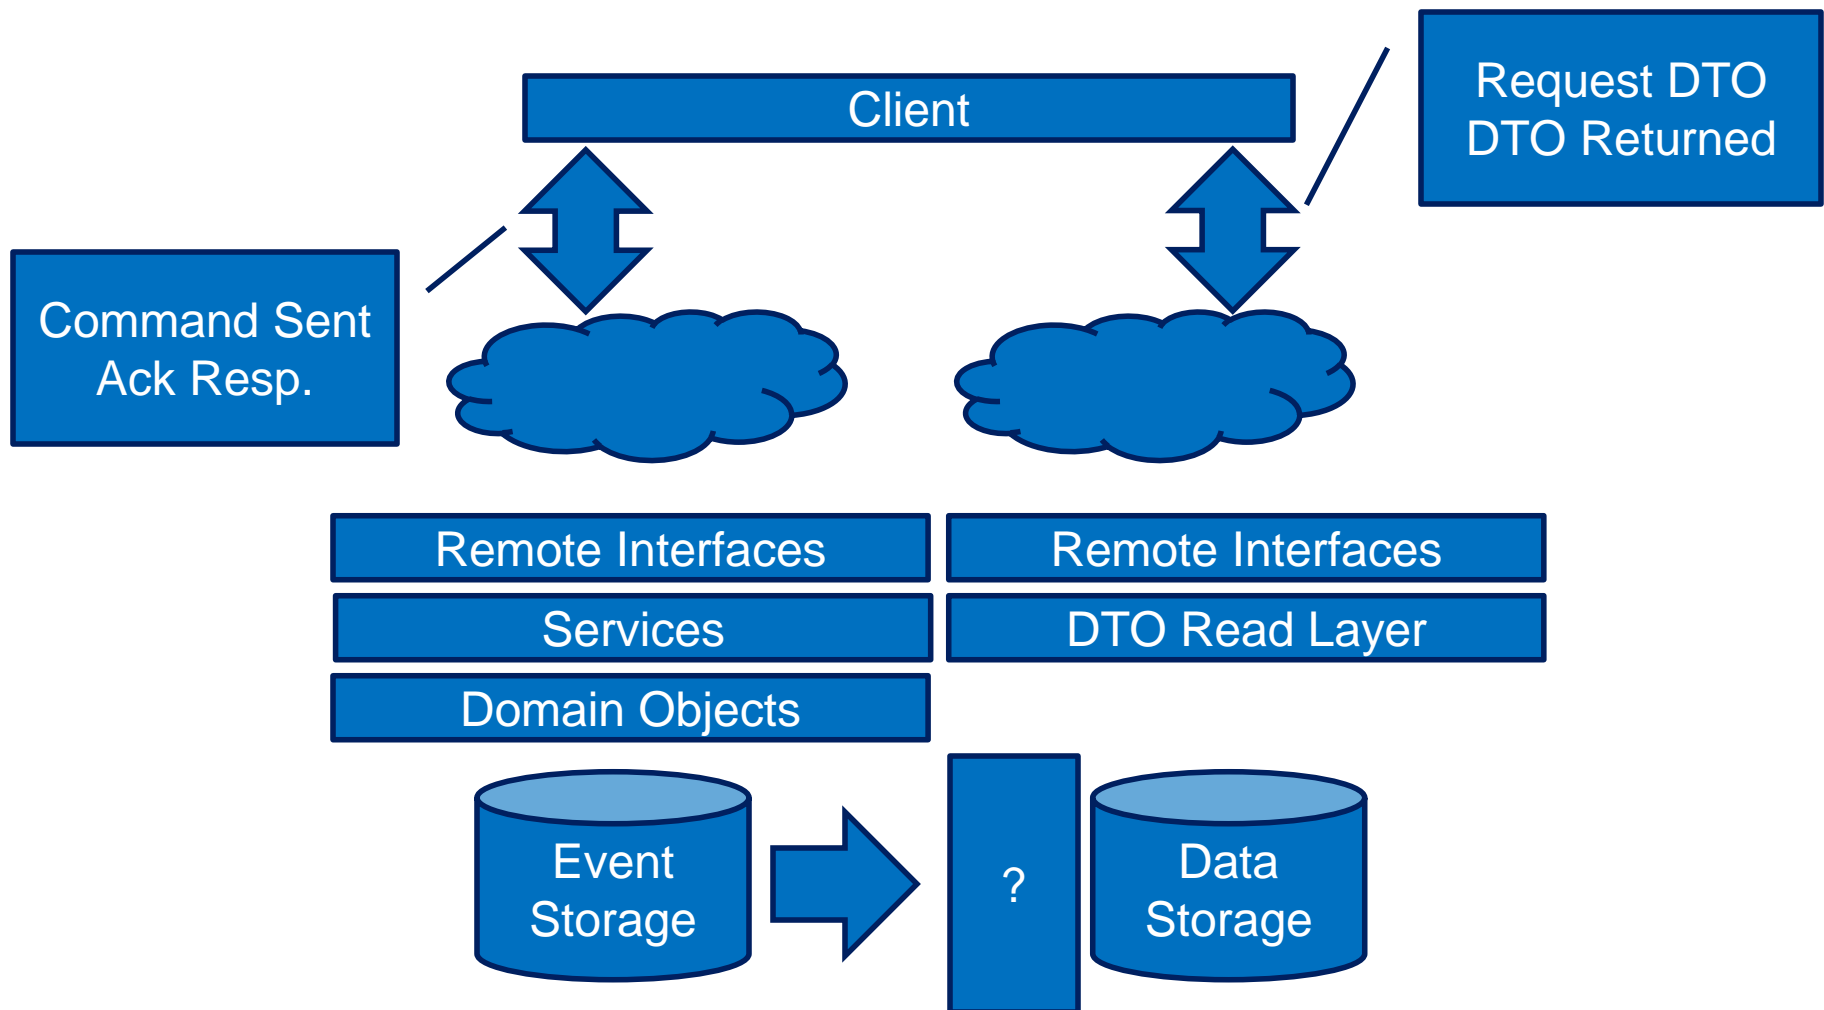
- Logstash + Elasticsearch + Kibana



# How to use Elasticsearch - CQRS



# How to use Elasticsearch - CQRS



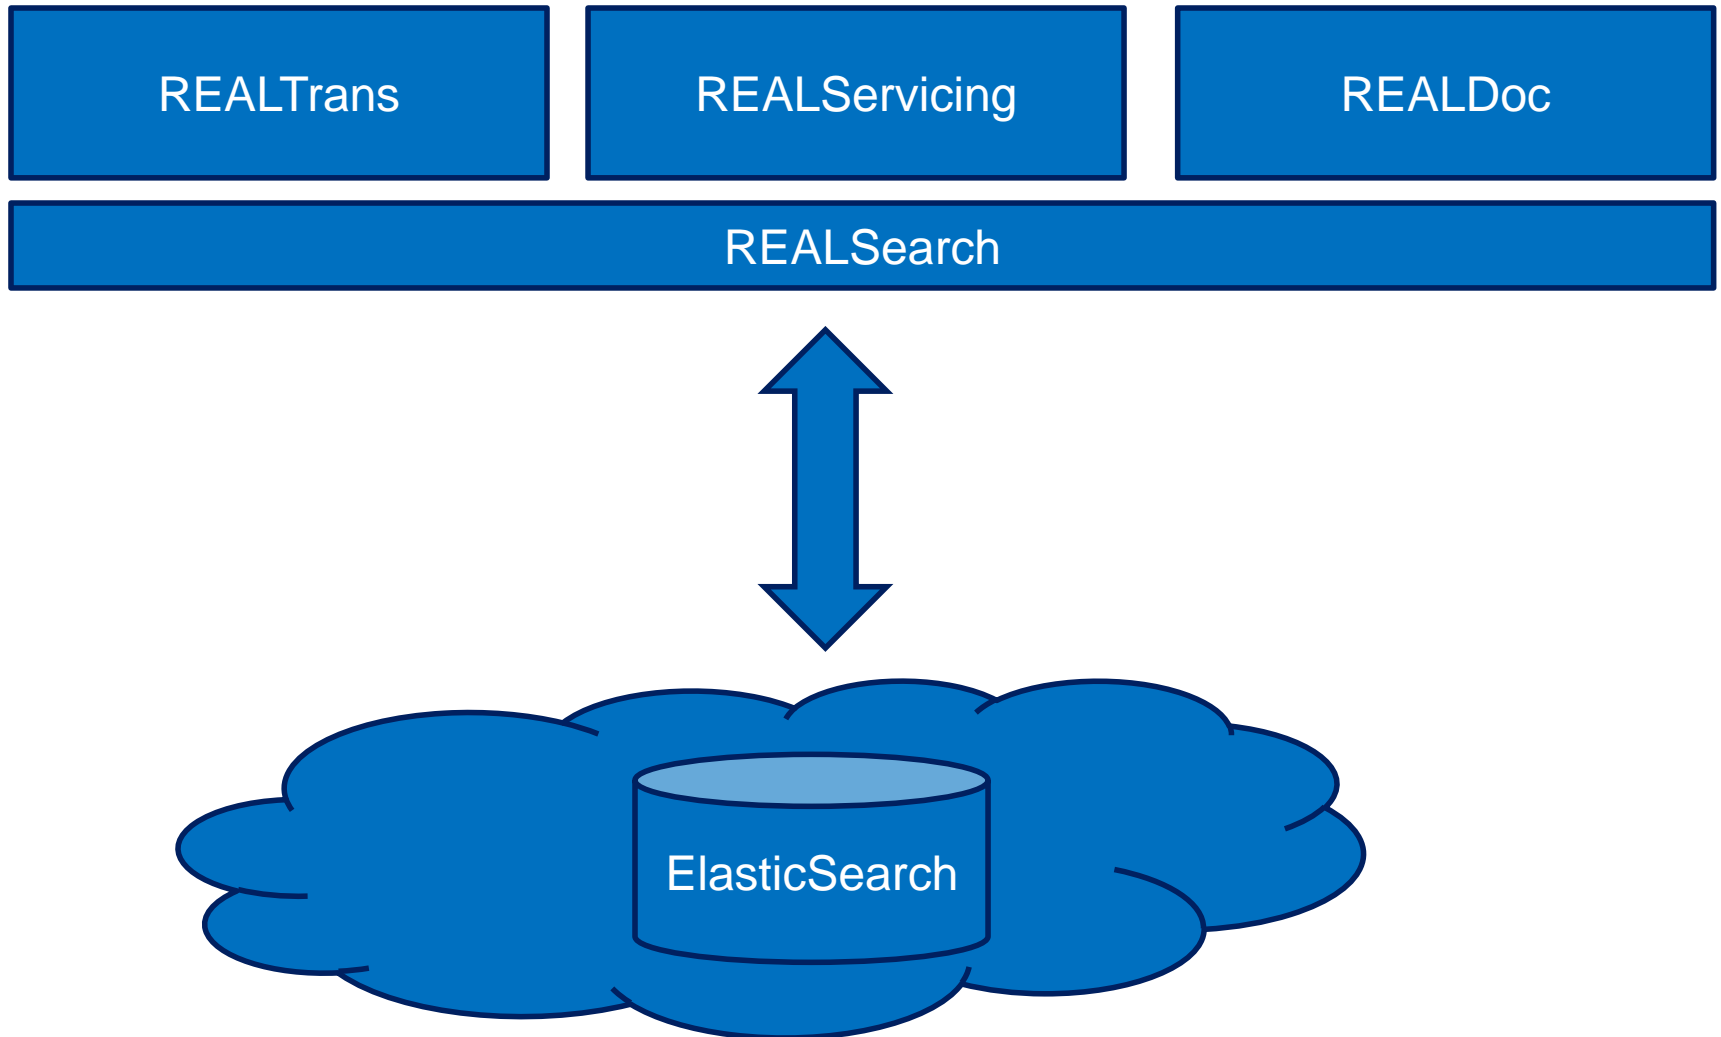


# A note on Rivers

- JDBC
- CouchDB
- MongoDB
- RabbitMQ
- Twitter
- And more...

```
"type" : "jdbc",  
"jdbc" : {  
    "driver" : "com.mysql.jdbc.Driver",  
    "url" : "jdbc:mysql://localhost:3306/my_db",  
    "user" : "root",  
    "password" : "mypassword",  
    "sql" : "select * from products"  
}
```

# ElasticSearch at Work



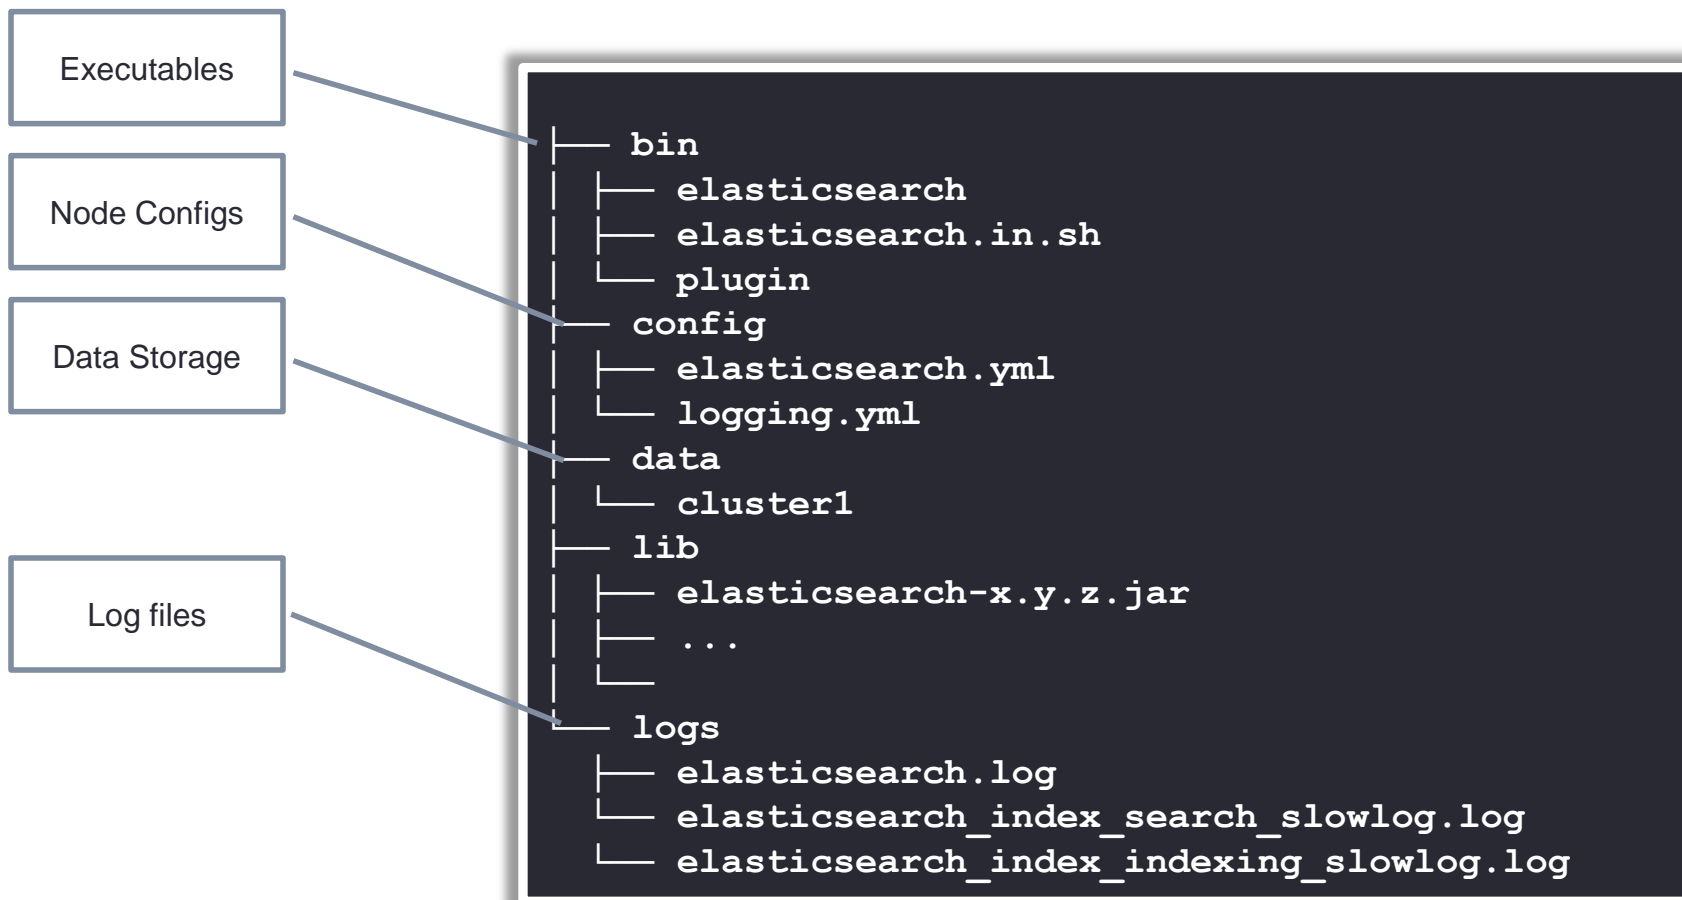
# What sucks about Elasticsearch

- No AUTH/AUTHZ
- No Usage Metrics

# How the World Uses Elasticsearch

# The Basics - Distro

- Download and Run



# The Basics - Glossary

- Node = One Elasticsearch instance (1 java proc)
- Cluster = 1..N Nodes w/ same Cluster Name
- Index = Similar to a DB
  - Named Collection of Documents
  - Maps to 1..N Primary shards && 0..N Replica shards
- Mapping Type = Similar to a DB Table
  - Document Definition
- Shard = One Lucene instance
  - Distributed across all nodes in the cluster.

# The Basics - Document Structure

- Modeled as a JSON object

```
{  
  "genre": "Crime",  
  "language": "English",  
  "country": "USA",  
  "runtime": 170,  
  "title": "Scarface",  
  "year": 1983  
}
```

```
{  
  "_index": "imdb",  
  "_type": "movie",  
  "_id": "u17o8zy9RcKg6SjQZqQ4Ow",  
  "_version": 1,  
  "_source": {  
    "genre": "Crime",  
    "language": "English",  
    "country": "USA",  
    "runtime": 170,  
    "title": "Scarface",  
    "year": 1983  
  }  
}
```

# The Basics - Document Structure

- Document Metadata fields
  - `_id`
  - `_type` : mapping type
  - `_source` : enabled/disabled
  - `_timestamp`
  - `_ttl`
  - `_size` : size of uncompressed `_source`
  - `_version`



# The Basics - Document Structure

- Mapping:
  - ES will auto-map (type) fields
  - You can specify mapping, if needed
- Data Types:
  - String
  - Number
    - Int, long, float, double, short, byte
  - Boolean
  - Datetime
    - formatted
  - geo\_point, geo\_shape
  - Array
  - Nested
  - IP

# A Mapping Type

```
"imdb": {  
  "movie": {  
    "properties": {  
      "country": {  
        "type": "string",  
        "store": true,  
        "index": false  
      },  
      "genre": {  
        "type": "string",  
        "null_value": "na",  
        "store": false,  
        "index": true  
      },  
      "year": {  
        "type": "long"  
      }  
    }  
  }  
}
```

# Lucene – Inverted Index

- Which presidential speeches contain the words “fair”
  - Go over every speech, word by word, and mark the speeches that contain it
  - Fails at large scale

# Lucene – Inverted Index

- Inverting
  - Take all the speeches
  - Break them down by word (tokenize)
  - For each word, store the IDs of the speeches
  - Sort all words (tokens)
- Searching
  - Finding the word is fast
  - Iterate over document IDs that are referenced

Token	Doc Frequency	Doc IDs
Jobs	2	4,8
Fair	5	1,2,4,8,42
Bush	300	1,2,3,4,5,6, ...

# Lucene – Inverted Index

- Not an algorithm
- Implementations vary


```
{0} - "Turtles love pizza"  
{1} - "I love my turtles"  
{2} - "My pizza is good"
```

	Record Level	Fully Inverted
"turtles"	{0, 1}	{ (0, 0), (1, 3) }
"love"	{0, 1}	{ (0, 1), (1, 1) }
"pizza"	{0, 2}	{ (0, 2), (2, 1) }
"i"	{1}	{ (1, 0) }
"my"	{1, 2}	{ (1, 2), (2, 0) }
"is"	{2}	{ (2, 2) }
"good"	{2}	{ (2, 3) }

```
"turtles"  {0, 1}  
"my"       {1, 2}
```

# Cluster Topology

- 4 Node Cluster
- Index Configuration:
  - “A”: 2 Shards, 1 Replica
  - “B”: 3 Shards, 1 Replica



# Building a Cluster

## Start Cluster...

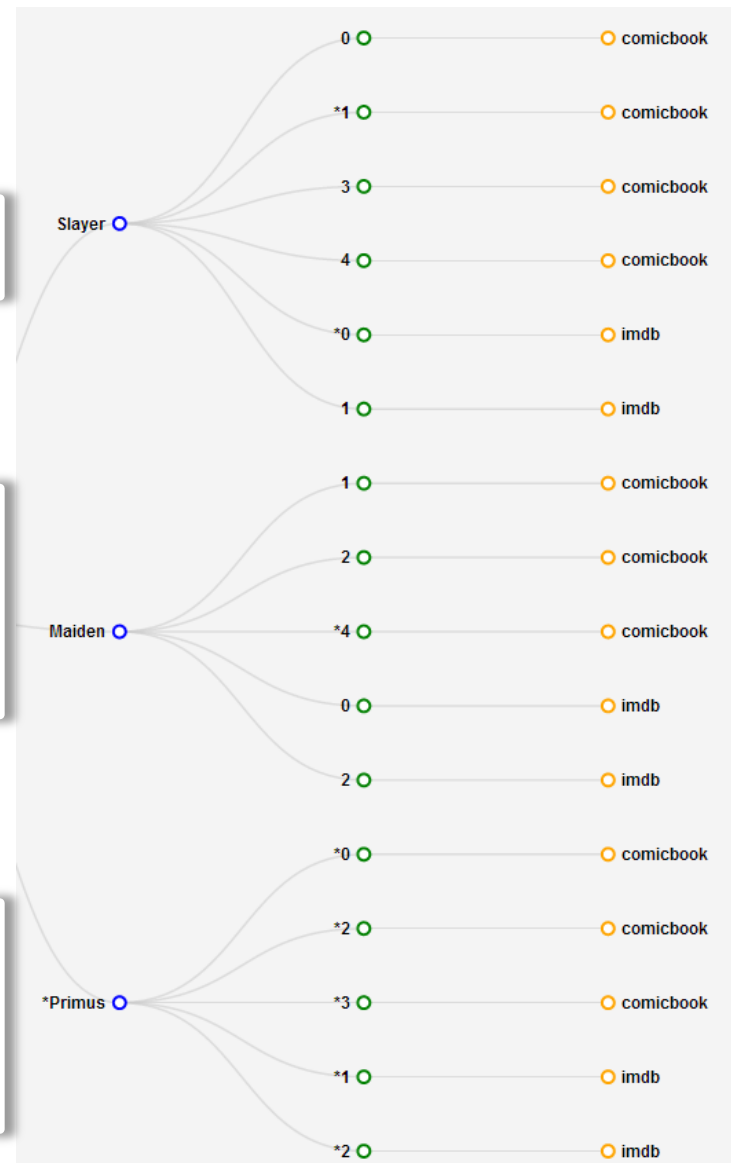
```
start cmd.exe /C elasticsearch -Des.node.name=Primus
start cmd.exe /C elasticsearch -Des.node.data=true -Des.node.master=false -Des.node.name=Slayer
start cmd.exe /C elasticsearch -Des.node.data=true -Des.node.master=true -Des.node.name=Maiden
```

## Create Index...

```
curl -XPUT 'http://localhost:9200/imdb' -d '{
  "settings": {
    "index": {
      "number_of_shards": 3,
      "number_of_replicas": 1
    }
  }
}'
```

## Index Document...

```
curl -XPOST 'http://localhost:9200/imdb/movie/' -d '{
  "genre": "Comedy",
  "language": "English",
  "country": "USA",
  "runtime": 99,
  "title": "Big Trouble in Little China",
  "year": 1986
}'
```



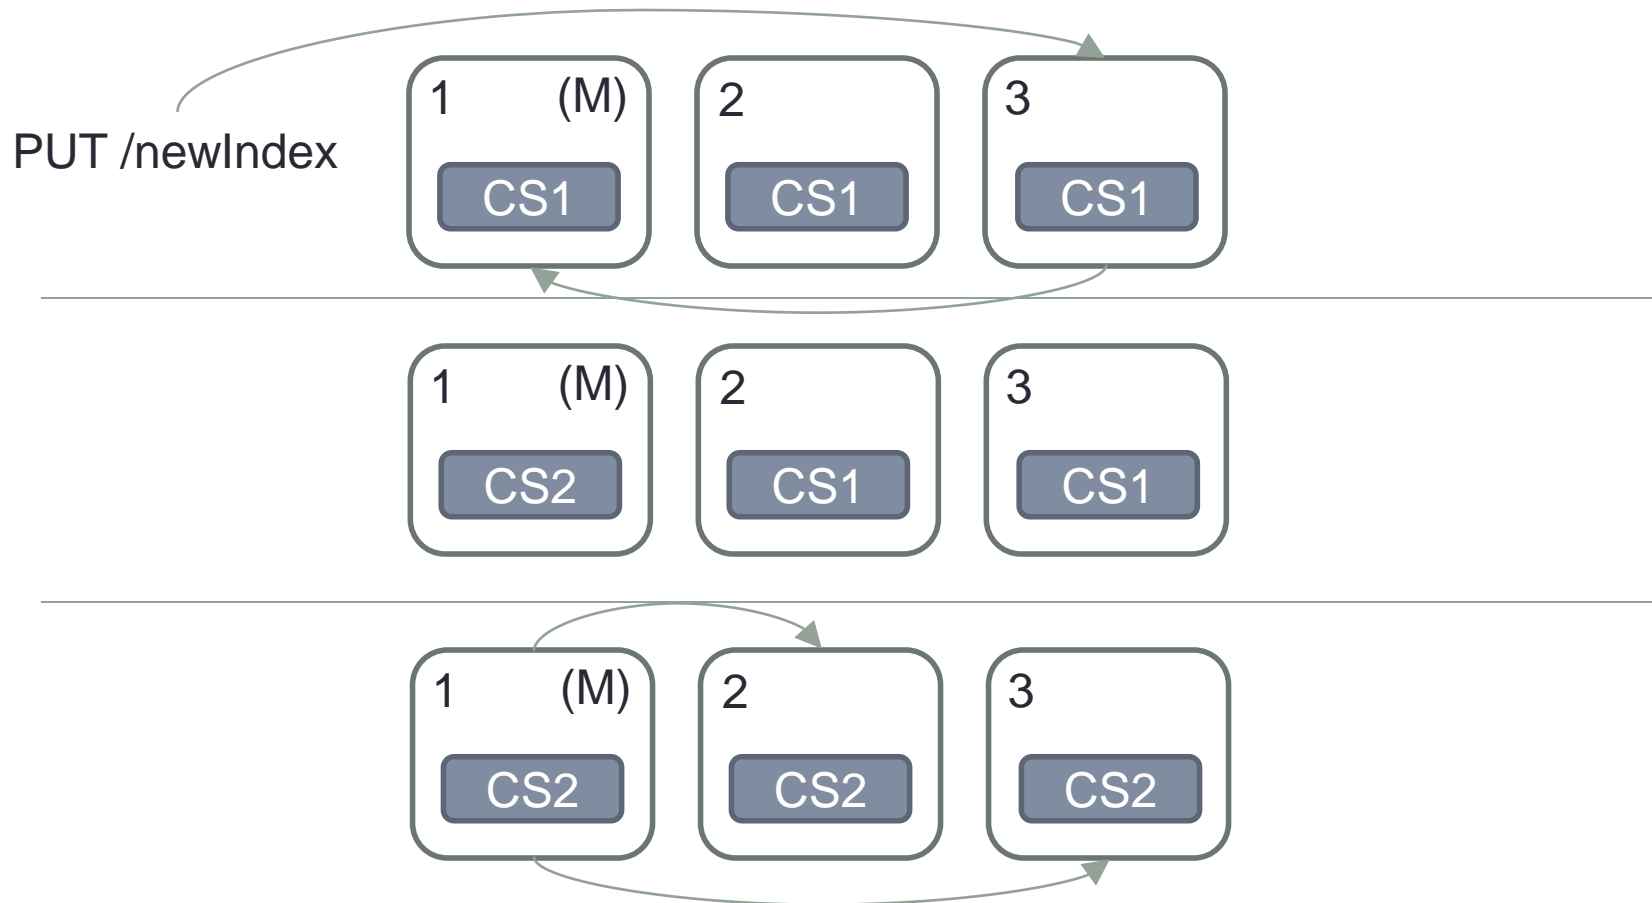
# Cluster State

- Cluster State
  - Node Membership
  - Indices Settings and Mappings (Types)
  - Shard Allocation Table
  - Shard State
- cURL -XGET `http://localhost:9200/_cluster/state?pretty=1'`



# Cluster State

- Changes in State published from Master to other nodes



# Discovery

- Nodes discover each other using multicast.
  - Unicast is an option

```
discovery.zen.ping.multicast.enabled: false  
  
discovery.zen.ping.unicast.hosts: ["host1", "host2:port", "host3"]
```


- Each cluster has an elected master node
  - Beware of split-brain

# The Basics - Shards

- Primary Shard:
  - First time Indexing
  - Index has 1..N primary shards (default: 5)
  - # Not changeable once index created
- Replica Shard:
  - Copy of the primary shard
  - Can be changed later
  - Each primary has 0..N replicas
  - HA:
    - Promoted to primary if primary fails
    - Get/Search handled by primary||replica

# Shard Auto-Allocation

- Add a node - Shards Relocate



- Shard Stages
  - UNASSIGNED
  - INITIALIZING
  - STARTED
  - RELOCATING

# The Basics – Searching

- How it works:
  - Search request hits a node
  - Node broadcasts to every shard in the index
  - Each shard performs query
  - Each shard returns results
  - Results merged, sorted, and returned to client.
- Problems:
  - ES has no idea where your document is
  - Broadcast query to 100 nodes
  - Performance degrades

# The Basics - Shards

- Shard Allocation Awareness
  - `cluster.routing.allocation.awareness.attributes: rack_id`
  - Example:
    - 2 Nodes with `node.rack_id=rack_one`
    - Create Index 5 shards / 1 replica (10 shards)
    - Add 2 Nodes with `node.rack_id=rack_two`
    - Shards RELOCATE to even distribution
    - Primary & Replica will NOT be on the same `rack_id` value.
- Shard Allocation Filtering
  - `node.tag=val1`
  - `index.routing.allocation.include.tag:val1,val2`

```
curl -XPUT localhost:9200/newIndex/_settings -d '{  
  "index.routing.allocation.include.tag" : "val1,val2"  
}'
```

# Nodes

- Master node handles cluster-wide (Meta-API) events:
  - Node participation
  - New indices create/delete
  - Re-Allocation of shards
- Data Nodes
  - Indexing / Searching operations
- Client Nodes
  - REST calls
  - Light-weight load balancers

# REST API

- Create Index
  - `action.auto_create_index: 0`
- Index Document
  - Dynamic type mapping
  - Versioning
  - ID specification
  - Parent / Child (`/1122?parent=1111`)



# REST API – Versioning

- Every document is Versioned
- Version assigned on creation
  - Version number can be assigned

# REST API - Update

- Update using partial data
- Partial doc merged with existing
- Fails if document doesn't exist
- “Upsert” data used to create a doc, if doesn't exist

```
{  
  "upsert" : {  
    "title": "Blade Runner"  
  }  
}
```

# REST API

- Exists
  - No overhead in loading
  - Status Code Result
- Delete
- Get
  - Multi-Get

```
{
  "docs" : [
    {
      "_id" : "1"
      "_index" : "imdb"
      "_type" : "movie"
    },
    {
      "_id" : "5"
      "_index" : "oldmovies"
      "_type" : "movie"
      "_fields" : ["title", "genre"]
    }
  ]
}
```

# REST API - Search

- Free Text Search
  - URL Request
    - [http://localhost:9200/imdb/movie/\\_search?q=scar\\*](http://localhost:9200/imdb/movie/_search?q=scar*)
- Complex Query
  - [http://localhost:9200/imdb/movie/\\_search?q=scarface+OR+star](http://localhost:9200/imdb/movie/_search?q=scarface+OR+star)
  - [http://localhost:9200/imdb/movie/\\_search?q=\(scarface+OR+star\)+AND+year:\[1981+TO+1984\]](http://localhost:9200/imdb/movie/_search?q=(scarface+OR+star)+AND+year:[1981+TO+1984])

# REST API - Search

- Search Types:
  - [http://localhost:9200/imdb/movie/\\_search?q=\(scarface+OR+star\)+AND+year:\[1941+TO+1984\]&search\\_type=count](http://localhost:9200/imdb/movie/_search?q=(scarface+OR+star)+AND+year:[1941+TO+1984]&search_type=count)
  - [http://localhost:9200/imdb/movie/\\_search?q=\(scarface+OR+star\)+AND+year:\[1941+TO+1984\]&search\\_type=query\\_then\\_fetch](http://localhost:9200/imdb/movie/_search?q=(scarface+OR+star)+AND+year:[1941+TO+1984]&search_type=query_then_fetch)
  - Query and Fetch (fastest):
    - Executes on all shards and return results
  - Query then Fetch (default):
    - Executes on all shards. Only some information returned for rank/sort, only the relevant shards are asked for data

# REST API – Query DSL

[http://localhost:9200/imdb/movie/\\_search?q=\(scarface+OR+star\)+AND+year:\[1981+TO+1984\]](http://localhost:9200/imdb/movie/_search?q=(scarface+OR+star)+AND+year:[1981+TO+1984])

Becomes...

```
curl -XPOST 'localhost:9200/_search?pretty' -d '{
  "query" : {
    "bool" : {
      "must" : [
        {
          "query_string" : {
            "query" : "scarface or star"
          }
        },
        {
          "range" : {
            "year" : { "gte" : 1931 }
          }
        }
      ]
    }
  }
}
```

# REST API – Query DSL

- Query String Request use Lucene query syntax
  - Limited
  - Instead use “match” query

```
curl -XPOST 'localhost:9200/_search?pretty' -d '{  
  "query" : {  
    "bool" : {  
      "must" : [  
        {  
          "match" : {  
            "message" : "scarface star"  
          }  
        },  
        {  
          "range" : {  
            "year" : { "gte" : 1981 }  
          }  
        }  
      ]  
    }  
  }  
}
```

...

Automatically builds  
a boolean query

# REST API – Query DSL

- Match Query

```
{
  "match":{
    "title":{
      "type":"phrase",
      "query":"quick fox",
      "slop":1
    }
  }
}
```

- Boolean Query

- Must: document must match query
- Must\_not: document must not match query
- Should: document doesn't have to match
  - If it matches... higher score

```
{
  "bool":{
    "must":[
      {
        "match":{
          "color":"blue"
        }
      },
      {
        "match":{
          "title":"shirt"
        }
      }
    ],
    "must_not":[
      {
        "match":{
          "size":"xxl"
        }
      }
    ],
    "should":[
      {
        "match":{
          "textile":"cotton"
        }
      }
    ]
  }
}
```



# REST API – Query DSL

- Range Query
  - Numeric / Date Types
- Prefix/Wildcard Query
  - Match on partial terms
- RegExp Query

```
{  
  "range":{  
    "founded_year":{  
      "gte":1990,  
      "lt":2000  
    }  
  }  
}
```

# REST API – Query DSL

- Geo\_bbox
  - Bounding box filter
- Geo\_distance
  - Geo\_distance\_range

```
{
  "query":{
    "filtered":{
      "query":{
        "match_all":{
          }
        },
      },
    "filter":{
      "geo_distance":{
        "distance":"400km"
        "location":{
          "lat":40.73,
          "lon":-74.1
        }
      }
    }
  }
}
```

```
{
  "query":{
    "filtered":{
      "query":{
        "match_all":{
          }
        },
      },
    "filter":{
      "geo_bbox":{
        "location":{
          "top_left":{
            "lat":40.73,
            "lon":-74.1
          },
          "bottom_right":{
            "lat":40.717,
            "lon":-73.99
          }
        }
      }
    }
  }
  ...
}
```

# REST API – Bulk Operations

- Bulk API
  - Minimize round trips with index/delete ops
  - Individual response for every request action
    - In order
  - Failure of one action will not stop subsequent actions.
  - localhost:9200/\_bulk

```
{ "delete" : { "_index" : "imdb", "_type" : "movie", "_id" : "2" } }\n{ "index" : { "_index" : "imdb", "_type" : "actor", "_id" : "1" } }\n{ "first_name" : "Tony", "last_name" : "Soprano" }\n...\n{ "update" : { "_index" : "imdb", "_type" : "movie", "_id" : "3" } }\n{ doc : { "title" : "Blade Runner" } }\n
```

# Percolate API

- Reversing Search
  - Store queries and filter (percolate) documents through them.
  - Useful for Alert/Monitoring systems

```
curl -XPUT localhost:9200/_percolator/stocks/alert-on-nokia -d '{
  "query" : {
    "boolean" : {
      "must" : [
        { "term" : { "company" : "NOK" } },
        { "range" : { "value" : { "lt" : "2.5" } } }
      ]
    }
  }
}'
```

```
curl -X PUT localhost:9200/stocks/stock/1?percolate=* -d '{
  "doc" : {
    "company" : "NOK",
    "value" : 2.4
  }
}'
```

# Clients

- Client list: <http://www.elasticsearch.org/guide/clients/>
  - Java Client, JS, PHP, Perl, Python, Ruby
- Spring Data:
  - Uses TransportClient
  - Implementation of ElasticsearchRepository aligns with generic Repository interfaces.
  - ElasticsearchCrudRepository extends PagingandSortingRepository
  - <https://github.com/spring-projects/spring-data-elasticsearch>

```
@Document(indexName = "book", type = "book", indexStoreType = "memory", shards = 1, replicas = 0, refreshInterval = "-1")
public class Book {
    ...
}

public interface ElasticSearchBookRepository extends ElasticsearchRepository<Book, String> {
}
```

# B'what about Mongo?

- Mongo:
  - General purpose DB
- ElasticSearch:
  - Distributed text search engine

... that's all I have to say about that.

# Questions?

